

# FaaSLight Serverless Simulation User Guide

Mohab

May 27, 2025

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Prerequisites</b>	<b>2</b>
<b>3</b>	<b>Installation</b>	<b>2</b>
<b>4</b>	<b>Usage</b>	<b>2</b>
<b>5</b>	<b>Architecture</b>	<b>3</b>
<b>6</b>	<b>Key Functions</b>	<b>3</b>
6.1	faaslight_simulator_.py . . . . .	3
6.2	app.py . . . . .	4
<b>7</b>	<b>Calling Sequence</b>	<b>5</b>
<b>8</b>	<b>Example Output</b>	<b>5</b>
<b>9</b>	<b>Troubleshooting</b>	<b>5</b>
<b>10</b>	<b>Two-Replica Setup for Improved Performance</b>	<b>6</b>

# 1 Introduction

FaaSLight is a Python-based project that simulates and compares serverless computing models, evaluating five platforms: FaaSLight Original, FaaSLight Enhanced, SAND++, AWS Fargate, and Unikernel. It measures performance across compute, I/O, and memory-intensive workloads, focusing on latency, CPU/memory usage, security, cost, and a composite score. The project uses Docker containers to host Flask APIs, a Python simulator to generate workloads, and Plotly for interactive visualizations. This guide provides instructions for setup, usage, and details on the project's architecture, functions, and execution flow.

## 2 Prerequisites

- **Operating System:** Windows (tested), Linux, or macOS.
- **Docker Desktop:** Latest version.
- **Python:** 3.10 or higher.
- **Disk Space:** Approximately 500MB for Docker images and results.
- **RAM:** 8GB recommended (4GB minimum).
- **CPU:** 4 cores recommended.

## 3 Installation

### 1. Clone the Repository:

```
1 git clone https://github.com/<your-username>/faaslight.git
2 cd faaslight
3
```

### 2. Set Up Python Virtual Environment:

```
1 python -m venv venv
2 source venv/Scripts/activate % Windows
3 % or source venv/bin/activate % Linux/macOS
4 pip install -r requirements.txt
5
```

### 3. Start Docker Desktop: Ensure Docker is running.

## 4 Usage

### 1. Start Docker Containers:

```
1 docker-compose up -d --build
2
```

Verify containers are running:

```
1 docker-compose ps
2
```

Expected output:

Name	Command	State
faaslight_project_new-faaslight_enhanced-1	flask run --host=0.0.0.0	Up
faaslight_project_new-faaslight_original-1	flask run --host=0.0.0.0	Up
...		

## 2. Run the Simulation:

```
1 python faaslight_simulator_.py
2
```

The simulation takes approximately 1–2 minutes, prints performance metrics, and generates `results/metrics.csv` and `results/dashboard.html`.

## 3. View Results: Open `results/dashboard.html` in a web browser to view interactive bar and scatter plots.

## 4. Clean Up:

```
1 docker-compose down
2 deactivate
3
```

# 5 Architecture

The FaaSLight project simulates serverless platforms using a Docker-based architecture:

- **Docker Containers:** Five services (`faaslight_original`, `faaslight_enhanced`, `sand_plus`, `fargate_inspired`, `unikernel_inspired`) are defined in `docker-compose.yml` and built from `Dockerfile` using Python 3.9-slim.
- **Flask APIs:** The `app.py` script runs a Flask server in each container, providing endpoints (`/compute`, `/io`, `/memory`, `/verify`) to simulate serverless workloads.
- **Simulator:** The `faaslight_simulator_.py` script sends HTTP requests to the containers, collects performance metrics, and generates visualizations using Plotly.
- **Network:** A bridge network (`faaslight-net`) connects containers to the host, mapping container port 5000 to localhost ports 32768–32772.
- **Output:** Results are saved as `results/metrics.csv` (raw metrics) and `results/dashboard.htm` (interactive plots).

# 6 Key Functions

The following are the primary functions in `faaslight_simulator_.py` and `app.py`:

## 6.1 `faaslight_simulator_.py`

- `send_request(url, func_type, is_cold=False):`

- *Purpose*: Sends an HTTP GET request to a service endpoint (e.g., `http://localhost:32769`).
- *Parameters*:
  - \* `url`: Service URL (e.g., `http://localhost:32769`).
  - \* `func_type`: Workload type (`compute`, `io`, `memory`).
  - \* `is_cold`: If `True`, applies a cold-start penalty (e.g., 10ms for FaaSLight Enhanced).
- *Returns*: Dictionary containing `latency` (ms), `status` (HTTP code), `type`, `security_overhead` (ms), `cpu_usage` (%), `memory_usage` (%).
- *Example*:

```

1 result = send_request("http://localhost:32769", "compute", is_cold
    =True)
2 # Returns: {"latency": 750, "status": 200, "type": "compute", "
    security_overhead": 10, ...}
3

```

- **`run_simulation()`**:
  - *Purpose*: Executes the simulation for all services, sending 200 requests per service with 10% cold starts.
  - *Logic*: Uses `ThreadPoolExecutor` with 10 workers to parallelize requests. Collects metrics and performs security verification.
  - *Returns*: Dictionary of raw metrics per service and workload.
- **`compute_metrics(metrics)`**:
  - *Purpose*: Processes raw metrics into average and P95 latency, CPU/memory usage, security score, cost, and composite score.
  - *Formulas*:
    - \* Cost:  $(avg\_cpu + avg\_memory) \times avg\_latency / 1000$ , reduced by 0.8 for FaaSLight Enhanced.
    - \* Composite:  $(security \times 1000) / (avg\_latency + p95\_latency + cost)$ .
  - *Returns*: Dictionary of processed metrics.
- **`plot_results(results)`**:
  - *Purpose*: Generates interactive visualizations using Plotly: a bar chart of composite scores and a scatter plot of latency vs. security.
  - *Output*: Saves `results/metrics.csv` and `results/dashboard.html`.

## 6.2 app.py

- **`compute()`**: Flask route simulating CPU-intensive tasks (e.g., matrix operations with NumPy).
- **`io()`**: Flask route simulating network-bound tasks (e.g., HTTP requests).

- **memory()**: Flask route simulating memory-intensive tasks (e.g., large array allocation).
- **verify()**: Flask route simulating a security check (e.g., SHA-256 hash), returning HTTP 200.

## 7 Calling Sequence

The execution flow of `faaslight_simulator_.py` is as follows:

1. **Main Block**: Initiates the simulation, printing the start timestamp.
2. **run\_simulation()**:
  - Iterates over services (e.g., FaaSLight Enhanced).
  - Sends 200 requests per service using `send_request()`, with 20 cold starts.
  - Sends a `/verify` request for security checks.
3. **compute\_metrics()**: Processes raw metrics into performance indicators.
4. **plot\_results()**: Saves results to `metrics.csv` and generates `dashboard.html`.
5. **Completion**: Prints a message indicating output file locations.

## 8 Example Output

Below is an example output from a simulation run (single replica, 2025-05-27 13:06:06):

```
Starting FaaSLight simulation at 2025-05-27 13:06:06
```

```
Testing FaaSLight Enhanced...
```

```
FaaSLight Enhanced (compute): Avg Latency ~749ms, P95 Latency ~1488ms, CPU ~13.5%, Me
```

```
FaaSLight Enhanced (io): Avg Latency ~564ms, P95 Latency ~1206ms, CPU ~12.2%, Memory
```

```
FaaSLight Enhanced (memory): Avg Latency ~1095ms, P95 Latency ~1602ms, CPU ~12.1%, Me
```

```
...
```

```
Simulation complete. Results saved to results/dashboard.html and results/metrics.csv
```

With two replicas (recommended):

```
FaaSLight Enhanced (compute): Avg Latency ~150ms, P95 Latency ~180ms, CPU ~20.0%, Mem
```

```
FaaSLight Enhanced (io): Avg Latency ~50ms, P95 Latency ~70ms, CPU ~10.0%, Memory ~15
```

## 9 Troubleshooting

- **Containers Not Running:**

```
1 docker-compose logs faaslight_project_new-faaslight-enhanced-1
2
```

- **Port Conflicts:**

```
1 netstat -aon | findstr "32768 32769"
2 taskkill /PID <pid> /F
3
```

- **High Latency:** Switch to two replicas (see Section 9).
- **Dependency Issues:**

```
1 pip install -r requirements.txt
2
```

## 10 Two-Replica Setup for Improved Performance

To achieve lower latencies ( 150ms compute, 4.20 I/O composite score), configure two replicas for `faaslight_enhanced`:

### 1. Edit `docker-compose.yml`:

```
1 faaslight_enhanced:
2   ...
3   expose:
4     - "5000"
5   deploy:
6     replicas: 2
7     resources:
8       limits:
9         cpus: '0.75'
10        memory: '768M'
11
```

Remove `container_name` and `ports`: `["32769:5000"]`.

### 2. Edit `faaslight_simulator.py`:

```
1 def get_service_ports(service_name, num_replicas=1, project_prefix="
faaslight_project_new"):
2     import subprocess
3     ports = []
4     for i in range(1, num_replicas + 1):
5         try:
6             container_name = f"{project_prefix}-{service_name}-{i}"
7             cmd = f"docker-compose port {container_name} 5000"
8             output = subprocess.check_output(cmd, shell=True, text=
True).strip()
9             if output:
10                 host_port = output.split(":")[1]
11                 ports.append(f"http://localhost:{host_port}")
12         except subprocess.CalledProcessError as e:
13             print(f"Error getting port for {container_name}: {e}")
14     return ports
15
16 services["FaaSLight Enhanced"] = {"url": None, "cold_start_penalty":
10, "scaling_factor": 2}
17 services["FaaSLight Enhanced"]["urls"] = get_service_ports("
faaslight_enhanced", num_replicas=2)
18 if not services["FaaSLight Enhanced"]["urls"]:
19     raise RuntimeError("No ports found for faaslight_enhanced replicas
")
20
```

Update `run_simulation()` to use `urls` (see repository).

3. **Ensure Resources:** Configure Docker Desktop with 4 CPUs, 8GB (Settings > Resources).

4. **Run:**

```
1 docker-compose up -d --build
2 python faaslight_simulator_.py
3
```